

Navigation

October 1, 2019

1 Navigation

You are welcome to use this coding environment to train your agent for the project. Follow the instructions below to get started!

1.0.1 1. Start the Environment

Run the next code cell to install a few packages. This line will take a few minutes to run!

```
In [1]: !pip install unityagents
        !pip -q install ./python
        !pip install numpy==1.13.3
        from unityagents import UnityEnvironment
        import numpy as np
        import random
        import torch
        from collections import deque
        import matplotlib.pyplot as plt
        from dqn_agent import Agent
        %matplotlib inline
```

Collecting unityagents

Downloading <https://files.pythonhosted.org/packages/82/42/c337c5ba34c72a2e01fab4abe113ba9f282c>
100% || 81kB 3.8MB/s ta 0:00:01

Collecting tensorflow==1.7.1 (from unityagents)

Downloading <https://files.pythonhosted.org/packages/66/83/35c3f53129dfc80d65ebbe07ef0575263c3c>
100% || 48.1MB 869kB/s eta 0:00:01

Requirement already satisfied: numpy>=1.11.0 in /opt/conda/lib/python3.6/site-packages (from unityagents)

Requirement already satisfied: pytest>=3.2.2 in /opt/conda/lib/python3.6/site-packages (from unityagents)

Requirement already satisfied: pyyaml in /opt/conda/lib/python3.6/site-packages (from unityagents)

Collecting jupyter (from unityagents)

Downloading <https://files.pythonhosted.org/packages/83/df/0f5dd132200728a86190397e1ea87cd76244>

Collecting grpcio==1.11.0 (from unityagents)

Downloading <https://files.pythonhosted.org/packages/c8/b8/00e703183b7ae5e02f161dafacdfa8edbd72>
100% || 8.8MB 4.8MB/s eta 0:00:01

Requirement already satisfied: matplotlib in /opt/conda/lib/python3.6/site-packages (from unityagents)

Requirement already satisfied: Pillow>=4.2.1 in /opt/conda/lib/python3.6/site-packages (from unityagents)

Collecting protobuf==3.5.2 (from unityagents)

 Downloading https://files.pythonhosted.org/packages/9b/0d/6e9cf54be86c742a375a0ca0addadb582c89/protobuf-3.5.2.tar.gz

 100% || 6.4MB 10.1MB/s ta 0:00:01

Collecting docopt (from unityagents)

 Downloading https://files.pythonhosted.org/packages/a2/55/8f8cab2afd404cf578136ef2cc5dfb50baa1/docopt-0.6.2.tar.gz

Collecting termcolor>=1.1.0 (from tensorflow==1.7.1->unityagents)

 Downloading https://files.pythonhosted.org/packages/8a/48/a76be51647d0eb9f10e2a4511bf3ffb8cc1e/termcolor-1.1.0.tar.gz

Collecting gast>=0.2.0 (from tensorflow==1.7.1->unityagents)

 Downloading https://files.pythonhosted.org/packages/1f/04/4e36c33f8eb5c5b6c622a1f4859352a6acca/gast-0.2.0.tar.gz

Collecting absl-py>=0.1.6 (from tensorflow==1.7.1->unityagents)

 Downloading https://files.pythonhosted.org/packages/3c/0d/7cbf64cac3f93617a2b6b079c0182e4a83a3/absl_py-0.1.6.tar.gz

 100% || 112kB 20.2MB/s ta 0:00:01

Collecting astor>=0.6.0 (from tensorflow==1.7.1->unityagents)

 Downloading https://files.pythonhosted.org/packages/d1/4f/950dfae467b384fc96bc6469de25d832534f/astor-0.6.0.tar.gz

Requirement already satisfied: six>=1.10.0 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: wheel>=0.26 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Collecting tensorboard<1.8.0,>=1.7.0 (from tensorflow==1.7.1->unityagents)

 Downloading https://files.pythonhosted.org/packages/0b/ec/65d4e8410038ca2a78c09034094403d23122/tensorboard-1.7.0.tar.gz

 100% || 3.1MB 11.1MB/s ta 0:00:01

Requirement already satisfied: more-itertools>=4.0.0; python_version > "2.7" in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: pluggy!=0.10,<1.0,>=0.9 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: py>=1.5.0 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: attrs>=17.4.0 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: atomicwrites>=1.0 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: setuptools in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: wcwidth in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: nbconvert in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: notebook in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Collecting jupyter-console (from tensorflow==1.7.1->unityagents)

 Downloading https://files.pythonhosted.org/packages/cb/ee/6374ae8c21b7d0847f9c3722dcdfac986b8e/jupyter_console-6.0.0.tar.gz

Requirement already satisfied: ipywidgets in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Collecting qtconsole (from tensorflow==1.7.1->unityagents)

 Downloading https://files.pythonhosted.org/packages/d0/fc/4936c60be8d56acedfefcc13cf9f3d881e40/qtconsole-4.3.0.tar.gz

 100% || 122kB 19.2MB/s ta 0:00:01

Requirement already satisfied: ipykernel in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: python-dateutil>=2.0 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: pytz in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: cycler>=0.10 in /opt/conda/lib/python3.6/site-packages/cycler-0.1.0.tar.gz

Requirement already satisfied: pyparsing!=2.0.4,!2.1.2,!2.1.6,>=2.0.1 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: bleach==1.5.0 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: markdown>=2.6.8 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: html5lib==0.999999 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: werkzeug>=0.11.10 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: mistune>=0.8.1 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: jinja2 in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

Requirement already satisfied: pygments in /opt/conda/lib/python3.6/site-packages (from tensorflow==1.7.1->unityagents)

```

Requirement already satisfied: traitlets>=4.2 in /opt/conda/lib/python3.6/site-packages (from nbco
Requirement already satisfied: jupyter_core in /opt/conda/lib/python3.6/site-packages (from nbco
Requirement already satisfied: nbformat>=4.4 in /opt/conda/lib/python3.6/site-packages (from nbco
Requirement already satisfied: entrypoints>=0.2.2 in /opt/conda/lib/python3.6/site-packages (fro
Requirement already satisfied: pandocfilters>=1.4.1 in /opt/conda/lib/python3.6/site-packages (f
Requirement already satisfied: testpath in /opt/conda/lib/python3.6/site-packages (from nbconver
Requirement already satisfied: defusedxml in /opt/conda/lib/python3.6/site-packages (from nbconv
Requirement already satisfied: tornado>=4 in /opt/conda/lib/python3.6/site-packages (from notebo
Requirement already satisfied: pyzmq>=17 in /opt/conda/lib/python3.6/site-packages (from notebo
Requirement already satisfied: ipython_genutils in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: jupyter_client>=5.2.0 in /opt/conda/lib/python3.6/site-packages (
Requirement already satisfied: Send2Trash in /opt/conda/lib/python3.6/site-packages (from notebo
Requirement already satisfied: terminado>=0.8.1 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: prometheus_client in /opt/conda/lib/python3.6/site-packages (from
Collecting prompt-toolkit<2.1.0,>=2.0.0 (from jupyter-console->jupyter->unityagents)
  Downloading https://files.pythonhosted.org/packages/f7/a7/9b1dd14ef45345f186ef69d175bdd2491c40
    100% || 337kB 20.1MB/s ta 0:00:01
Requirement already satisfied: ipython in /opt/conda/lib/python3.6/site-packages (from jupyter-c
Collecting widgetsnbextension~=3.0.0 (from ipywidgets->jupyter->unityagents)
  Downloading https://files.pythonhosted.org/packages/8d/f2/c8bccccbed39d51d3e237fb0c0f0c9bbc84
    100% || 2.2MB 16.6MB/s ta 0:00:01
Requirement already satisfied: MarkupSafe>=0.23 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: decorator in /opt/conda/lib/python3.6/site-packages (from traitle
Requirement already satisfied: jsonschema!=2.5.0,>=2.4 in /opt/conda/lib/python3.6/site-packages
Requirement already satisfied: backcall in /opt/conda/lib/python3.6/site-packages (from ipython-
Requirement already satisfied: pickleshare in /opt/conda/lib/python3.6/site-packages (from ipyth
Requirement already satisfied: pexpect; sys_platform != "win32" in /opt/conda/lib/python3.6/site
Requirement already satisfied: simplegeneric>0.8 in /opt/conda/lib/python3.6/site-packages (from
Requirement already satisfied: jedi>=0.10 in /opt/conda/lib/python3.6/site-packages (from ipytho
Requirement already satisfied: ptyprocess>=0.5 in /opt/conda/lib/python3.6/site-packages (from p
Building wheels for collected packages: docopt, termcolor, gast, absl-py
  Running setup.py bdist_wheel for docopt ... done
  Stored in directory: /root/.cache/pip/wheels/9b/04/dd/7daf4150b6d9b12949298737de9431a324d4b797
  Running setup.py bdist_wheel for termcolor ... done
  Stored in directory: /root/.cache/pip/wheels/7c/06/54/bc84598ba1daf8f970247f550b175aaee85f68b
  Running setup.py bdist_wheel for gast ... done
  Stored in directory: /root/.cache/pip/wheels/59/38/c6/234dc39b4f6951a0768fbc02d5b7207137a5b1d9
  Running setup.py bdist_wheel for absl-py ... done
  Stored in directory: /root/.cache/pip/wheels/9a/1e/7a/456008eb5e47fd5de792c6139df6d5b3d5f71d51
Successfully built docopt termcolor gast absl-py
tensorflow 1.7.1 has requirement numpy>=1.13.3, but you'll have numpy 1.12.1 which is incompatib
ipython 6.5.0 has requirement prompt-toolkit<2.0.0,>=1.0.15, but you'll have prompt-toolkit 2.0.
Installing collected packages: protobuf, grpcio, termcolor, gast, absl-py, astor, tensorboard, t
  Found existing installation: protobuf 3.5.1
    Uninstalling protobuf-3.5.1:
      Successfully uninstalled protobuf-3.5.1
  Found existing installation: tensorflow 1.3.0
    Uninstalling tensorflow-1.3.0:

```

```

    Successfully uninstalled tensorflow-1.3.0
Found existing installation: prompt-toolkit 1.0.15
Uninstalling prompt-toolkit-1.0.15:
    Successfully uninstalled prompt-toolkit-1.0.15
Found existing installation: widgetsnbextension 3.1.0
Uninstalling widgetsnbextension-3.1.0:
    Successfully uninstalled widgetsnbextension-3.1.0
Successfully installed absl-py-0.8.0 astor-0.8.0 docopt-0.6.2 gast-0.3.2 grpcio-1.11.0 jupyter-1
Collecting numpy==1.13.3
  Downloading https://files.pythonhosted.org/packages/57/a7/e3e6bd9d595125e1abbe162e323fd2d06f6f
100% || 17.0MB 2.5MB/s eta 0:00:01
Installing collected packages: numpy
Found existing installation: numpy 1.12.1
Uninstalling numpy-1.12.1:
    Successfully uninstalled numpy-1.12.1
Successfully installed numpy-1.13.3

```

The environment is already saved in the Workspace and can be accessed at the file path provided below. Please run the next code cell without making any changes.

```

In [2]: # please do not modify the line below
        env = UnityEnvironment(file_name="/data/Banana_Linux_NoVis/Banana.x86_64")

```

```

INFO:unityagents:
'Academy' started successfully!
Unity Academy name: Academy
    Number of Brains: 1
    Number of External Brains : 1
    Lesson number : 0
    Reset Parameters :

Unity brain name: BananaBrain
    Number of Visual Observations (per agent): 0
    Vector Observation space type: continuous
    Vector Observation space size (per agent): 37
    Number of stacked Vector Observation: 1
    Vector Action space type: discrete
    Vector Action space size (per agent): 4
    Vector Action descriptions: , , ,

```

Environments contain *brains* which are responsible for deciding the actions of their associated agents. Here we check for the first brain available, and set it as the default brain we will be controlling from Python.

```

In [3]: # get the default brain
        brain_name = env.brain_names[0]
        # print(brain_name)

```

```

# academy_name = env.academy_name[1]
# print(academy_name)
brain = env.brains[brain_name]
# print(brain)

```

1.0.2 2. Examine the State and Action Spaces

Run the code cell below to print some information about the environment.

```

In [4]: # reset the environment
env_info = env.reset(train_mode=True)[brain_name]

# number of agents in the environment
print('Number of agents:', len(env_info.agents))

# number of actions
action_size = brain.vector_action_space_size
print('Number of actions:', action_size)

# examine the state space
state = env_info.vector_observations[0]
print('States look like:', state)
state_size = len(state)
print('States have length:', state_size)

```

```

Number of agents: 1
Number of actions: 4
States look like: [ 1.          0.          0.          0.          0.84408134  0.          0.
 1.          0.          0.0748472  0.          1.          0.          0.
 0.25755      1.          0.          0.          0.          0.74177343
 0.          1.          0.          0.          0.25854847  0.          0.
 1.          0.          0.09355672  0.          1.          0.          0.
 0.31969345  0.          0.          ]
States have length: 37

```

1.0.3 3. Take Random Actions in the Environment

In the next code cell, you will learn how to use the Python API to control the agent and receive feedback from the environment.

Note that **in this coding environment, you will not be able to watch the agent while it is training**, and you should set `train_mode=True` to restart the environment.

```

In [ ]: # env_info = env.reset(train_mode=True)[brain_name] # reset the environment
# state = env_info.vector_observations[0] # get the current state
# score = 0 # initialize the score
# while True:
#     action = np.random.randint(action_size) # select an action
#     env_info = env.step(action)[brain_name] # send the action to the environment

```

```

#     next_state = env_info.vector_observations[0]    # get the next state
#     reward = env_info.rewards[0]                  # get the reward
#     done = env_info.local_done[0]                  # see if episode has finished
#     score += reward                                # update the score
#     state = next_state                             # roll over the state to next time
#     if done:                                       # exit loop if episode finished
#         break

# print("Score: {}".format(score))

```

When finished, you can close the environment.

```
In [ ]: # env.close()
```

1.0.4 4. It's Your Turn!

Now it's your turn to train your own agent to solve the environment! A few **important notes**: - When training the environment, set `train_mode=True`, so that the line for resetting the environment looks like the following:

```
env_info = env.reset(train_mode=True)[brain_name]
```

- To structure your work, you're welcome to work directly in this Jupyter notebook, or you might like to start over with a new file! You can see the list of files in the workspace by clicking on *Jupyter* in the top left corner of the notebook.
- In this coding environment, you will not be able to watch the agent while it is training. However, *after training the agent*, you can download the saved model weights to watch the agent on your own machine!

```
In [5]: def dqn(n_episodes=5000, max_t=10000, eps_start=1.0, eps_end=0.01, eps_decay=0.97):
    scores = []                                # list containing scores from each episode
    scores_window = deque(maxlen=100)          # last 100 scores
    eps = eps_start                             # initialize epsilon
    for i_episode in range(1, n_episodes+1):
        env_info = env.reset(train_mode=True)[brain_name]
        state = env_info.vector_observations[0]
        score = 0
        for t in range(max_t):
            action = agent.act(state, eps)
            env_info = env.step(action)[brain_name]    # send the action to the env
            next_state = env_info.vector_observations[0] # get the next state
            reward = env_info.rewards[0]               # get the reward
            done = env_info.local_done[0]

            agent.step(state, action, reward, next_state, done)
            state = next_state
            score += reward
        if done:
            break

```

```

scores_window.append(score)          # save most recent score
scores.append(score)                 # save most recent score
eps = max(eps_end, eps_decay*eps)   # decrease epsilon
print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
if i_episode % 100 == 0:
    print('\rEpisode {} \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
if np.mean(scores_window) >= 13.0:
    print('\nEnvironment solved in {:d} episodes! \tAverage Score: {:.2f}'.format(i_episode, np.mean(scores_window)))
    torch.save(agent.qnetwork_local.state_dict(), 'checkpoint.pth')
    break
return scores

agent = Agent(state_size=37, action_size=4, seed=0)
scores = dqn()


# plot the scores
fig = plt.figure()
ax = fig.add_subplot(111)
plt.plot(np.arange(len(scores)), scores)
plt.ylabel('Score')
plt.xlabel('Episode #')
plt.show()

```

```

Episode 100      Average Score: 4.76
Episode 200      Average Score: 9.51
Episode 264      Average Score: 13.00
Environment solved in 164 episodes!      Average Score: 13.00

```



```
In [ ]:
```